

1. Q: What is a Database Management System (DBMS)? A: A DBMS is software used to manage and organize databases, enabling users to store, retrieve, update, and delete data efficiently.
2. Q: What are the advantages of using a DBMS? A: The advantages include data integrity, data security, data sharing, reduced data redundancy, and improved data retrieval and manipulation.
3. Q: What is data redundancy in a database? A: Data redundancy refers to the duplication of data in a database, which can lead to increased storage requirements and data inconsistency.
4. Q: What is a primary key in a database table? A: A primary key is a unique identifier for each record in a table, ensuring that each row can be uniquely identified.
5. Q: What is normalization in the context of databases? A: Normalization is the process of organizing data in a database to minimize data redundancy and dependency.
6. Q: What are the different levels of normalization? A: The levels of normalization are First Normal Form (1NF), Second Normal Form (2NF), Third Normal Form (3NF), and so on.
7. Q: What is a foreign key in a database table? A: A foreign key is a column or set of columns in a table that refers to the primary key of another table, creating a relationship between the two tables.
8. Q: What is the purpose of an index in a database? A: An index is used to speed up data retrieval operations by providing a quick access path to specific data within a table.
9. Q: What is a transaction in the context of databases? A: A transaction is a sequence of one or more database operations that are executed as a single unit of work, either all succeeding or all failing.
10. Q: What is ACID in database transactions? A: ACID stands for Atomicity, Consistency, Isolation, and Durability. It ensures that database transactions are reliable and maintain data integrity.
11. Q: What are the different types of database models? A: The main database models include the hierarchical model, network model, relational model, and object-oriented model.
12. Q: What is the difference between a database and a database management system? A: A database is a collection of structured data, while a database management system (DBMS) is software that manages and manipulates that data.
13. Q: What is the difference between a DBMS and a file system? A: A DBMS provides better data management, data integrity, and data security than a file system, which is primarily used for file storage and retrieval.
14. Q: What is a join operation in SQL? A: A join operation combines rows from two or more tables based on a related column between them.
15. Q: Explain the difference between INNER JOIN and OUTER JOIN in SQL. A: INNER JOIN returns only the matching rows from both tables, while OUTER JOIN returns matching rows and non-matching rows from one or both tables.
16. Q: What is a trigger in a database? A: A trigger is a stored procedure that automatically executes when a specific event occurs in the database, such as an update, insert, or delete operation.
17. Q: What is a database schema? A: A database schema is a blueprint that defines the structure of a database, including tables, columns, relationships, and constraints.
18. Q: What is data warehousing? A: Data warehousing is the process of collecting, storing, and organizing data from different sources to support business intelligence and reporting.

19. Q: What is OLAP (Online Analytical Processing)? A: OLAP is a category of software tools used to analyze and present multi-dimensional data for decision-making and reporting purposes.
20. Q: What is NoSQL database? A: NoSQL (Not Only SQL) databases are non-relational databases designed to handle large volumes of unstructured or semi-structured data.
21. Q: What is data mining in the context of databases? A: Data mining is the process of discovering patterns, relationships, and useful information from large datasets using various techniques and algorithms.
22. Q: What is a database view? A: A database view is a virtual table derived from one or more base tables in the database, presenting a filtered or summarized subset of data.
23. Q: What is data integrity in a database? A: Data integrity refers to the accuracy, consistency, and reliability of data stored in a database.
24. Q: What is a transaction log in a database? A: A transaction log is a file that records all the changes made to the database, ensuring that transactions can be rolled back or recovered in case of failures.
25. Q: What is the purpose of the COMMIT statement in SQL? A: The COMMIT statement is used to save all the changes made during a transaction permanently into the database.
26. Q: What is the purpose of the ROLLBACK statement in SQL? A: The ROLLBACK statement is used to undo all the changes made during a transaction and revert the database to its state before the transaction began.
27. Q: What is the difference between a clustered and a non-clustered index? A: A clustered index determines the physical order of data in a table, whereas a non-clustered index creates a separate structure to store the index.
28. Q: What is a data dictionary in a database? A: A data dictionary is a repository that stores metadata about the structure and contents of a database.
29. Q: How does a database handle concurrent transactions? A: A database uses locking mechanisms and isolation levels to ensure that multiple transactions can execute simultaneously without interfering with each other.
30. Q: What is a stored procedure in a database? A: A stored procedure is a precompiled set of SQL statements stored in the database, which can be executed as a single unit.
31. Q: What is a stored function in a database? A: A stored function is similar to a stored procedure but returns a value, making it useful for calculations or data manipulations.
32. Q: What is database replication? A: Database replication is the process of copying and maintaining a database on multiple servers to ensure data redundancy and availability.
33. Q: What is the difference between horizontal and vertical partitioning? A: Horizontal partitioning divides a table into multiple smaller tables with the same columns, while vertical partitioning separates columns into different tables.
34. Q: What is data normalization and denormalization? A: Data normalization is the process of reducing data redundancy and improving data integrity, while denormalization involves reintroducing redundancy to improve performance.
35. Q: What are composite keys in a database table? A: Composite keys are primary keys that consist of multiple columns used together to uniquely identify a row in a table.
36. Q: What is an ER diagram in database design? A: An Entity-Relationship (ER) diagram is a visual representation of the database's structure, showing entities, attributes, and relationships between them.

37. Q: What are surrogate keys, and why are they used? A: Surrogate keys are artificially created unique identifiers used as primary keys in a table to avoid reliance on natural keys, which may change or be ambiguous.
38. Q: What is the difference between a heap table and a clustered table? A: A heap table is an unordered collection of rows, while a clustered table has its rows stored in a specific order based on the clustered index.
39. Q: What is ACID compliance in databases? A: ACID compliance ensures that database transactions are Atomic, Consistent, Isolated, and Durable, providing reliability and consistency to the data.
40. Q: What is the purpose of the GROUP BY clause in SQL? A: The GROUP BY clause is used to group rows based on specified columns and allows performing aggregate functions on the grouped data.
41. Q: What is the role of a database administrator (DBA)? A: A database administrator is responsible for managing and maintaining the database system, including performance tuning, security, and backup and recovery.
42. Q: Explain the differences between SQL and NoSQL databases. A: SQL databases are relational and use structured query language, while NoSQL databases are non-relational and use various query languages like JSON, BSON, etc.
43. Q: What is the CAP theorem in database systems? A: The CAP theorem states that a distributed database can have at most two of the following three properties: Consistency, Availability, and Partition Tolerance.
44. Q: What is a data warehouse schema, and what are its types? A: A data warehouse schema determines how data is organized and stored. Types include Star schema, Snowflake schema, and Galaxy schema.
45. Q: What is the purpose of database constraints? A: Database constraints enforce rules and restrictions on data to maintain data integrity and consistency.
46. Q: What is an ETL process in data warehousing? A: ETL (Extract, Transform, Load) is a process that extracts data from various sources, transforms it to match the data warehouse schema, and loads it into the data warehouse.
47. Q: How does indexing impact database performance? A: Indexing improves data retrieval speed but may slow down data insertion and updates due to the overhead of maintaining the index.
48. Q: What is a data mart in data warehousing? A: A data mart is a subset of a data warehouse, focusing on specific business needs or departments.
49. Q: What is the difference between a clustered and non-clustered index? A: A clustered index determines the physical order of data in a table, while a non-clustered index creates a separate structure for fast data retrieval.
50. Q: What is the purpose of a database trigger? A: Database triggers are used to automatically execute predefined actions in response to specific events like data modification.
51. Q: How does a full-text search work in a database? A: A full-text search allows users to search for words or phrases within text columns of a database efficiently.
52. Q: What is the purpose of the SQL SELECT statement? A: The SQL SELECT statement retrieves data from one or more tables in a database.
53. Q: What is a self-join in SQL? A: A self-join is a query where a table is joined with itself to combine related rows.

54. Q: What is the purpose of the HAVING clause in SQL? A: The HAVING clause filters the results of aggregate functions applied to grouped data in SQL.
55. Q: What are the advantages and disadvantages of using database views? A: Advantages include data abstraction and security, while disadvantages involve performance overhead and complexity.
56. Q: What is database sharding? A: Database sharding is a technique of partitioning a large database into smaller, more manageable pieces across multiple servers.
57. Q: What is a materialized view in a database? A: A materialized view is a precomputed snapshot of a query result, stored for faster data retrieval.
58. Q: How can you improve the performance of database queries? A: Performance can be improved by using indexes, optimizing queries, caching, and denormalization.
59. Q: What is a database schema migration? A: Database schema migration is the process of updating a database's structure while preserving existing data.
60. Q: What is the role of a data analyst in database management? A: A data analyst analyzes data stored in databases to extract valuable insights for decision-making.
61. Q: What is a deadlock in a database, and how can it be resolved? A: A deadlock occurs when two or more transactions are unable to proceed because they are each waiting for a resource that is locked by another transaction. Deadlocks can be resolved by using techniques such as deadlock detection and timeout mechanisms.
62. Q: What is the purpose of the CASCADE option in foreign keys? A: The CASCADE option ensures that when a record in the referenced table is updated or deleted, all related records in the referencing table are automatically updated or deleted.
63. Q: What is a database index, and how does it work? A: A database index is a data structure that improves the speed of data retrieval operations by providing quick access to specific data within a table. It works similar to an index in a book, allowing faster lookup of specific data.
64. Q: What is a natural key in the context of databases? A: A natural key is a column or set of columns in a table that already exists in the real world and can be used as a unique identifier for the records.
65. Q: What is the purpose of the CHECK constraint in a database? A: The CHECK constraint restricts the values that can be inserted or updated in a column based on a specified condition.
66. Q: What are ACID properties, and why are they essential for a database? A: ACID properties (Atomicity, Consistency, Isolation, Durability) ensure that database transactions are reliable, maintain data integrity, and preserve consistency.
67. Q: What is a self-referential table in a database? A: A self-referential table is a table in which a foreign key references the primary key of the same table, establishing a relationship within the table itself.
68. Q: What is a database index fragmentation, and how can it be resolved? A: Index fragmentation occurs when data pages are out of order, leading to reduced query performance. It can be resolved by rebuilding or reorganizing indexes.
69. Q: What is the difference between a left join and a right join in SQL? A: A left join returns all rows from the left table and matching rows from the right table, while a right join returns all rows from the right table and matching rows from the left table.

70. Q: How can you ensure data security in a database system? A: Data security can be ensured by implementing access controls, encryption, strong passwords, and regularly auditing and monitoring the database.
71. Q: What is database replication latency, and how can it affect data consistency? A: Database replication latency refers to the delay in replicating data to secondary servers. High latency can result in data inconsistencies between the primary and secondary servers.
72. Q: What is a foreign key constraint in a database? A: A foreign key constraint ensures that the values in a column of one table correspond to the values of a primary key in another table, establishing a relationship between the two tables.
73. Q: What are the different types of database backups, and when should they be used? A: Types of database backups include full backups (complete backup of the entire database), differential backups (backup of changes since the last full backup), and incremental backups (backup of changes since the last backup, whether full or incremental). Each type has specific use cases and recovery implications.
74. Q: What is the purpose of the DML (Data Manipulation Language) in SQL? A: The DML is used to interact with the data in the database, allowing users to insert, update, delete, and retrieve data from tables.
75. Q: What is a surrogate data key, and how is it different from a natural data key? A: A surrogate data key is an artificial unique identifier, usually an auto-incremented number, used as a primary key. In contrast, a natural data key is based on real-world data and may not always be unique or stable.
76. Q: What is a B-tree index in a database? A: A B-tree index is a balanced tree structure used to organize data for efficient data retrieval in a database.
77. Q: What is a data dictionary, and why is it important in database management? A: A data dictionary is a repository that stores metadata about the database's structure and contents. It is crucial for data governance, documentation, and maintaining data consistency.
78. Q: What is a multi-version concurrency control (MVCC) in databases? A: MVCC is a technique used to manage concurrent access to data by allowing multiple versions of the same data to coexist during transactions, improving database performance and concurrency.
79. Q: How does a database handle data recovery after a system crash? A: Databases use transaction logs and other mechanisms to recover data to a consistent state after a system crash or failure.
80. Q: What is a database connection pool, and why is it used? A: A database connection pool is a cache of established database connections that can be reused, reducing the overhead of establishing new connections and improving performance.
81. Q: What is data compression in a database, and how does it benefit storage? A: Data compression is the process of reducing the size of data to save storage space, leading to lower storage costs and improved data retrieval speed.
82. Q: What is the difference between a clustered and non-clustered index in terms of data sorting? A: A clustered index determines the physical order of data in a table, whereas a non-clustered index creates a separate structure to map the index to the data.
83. Q: What are materialized views, and how do they differ from regular views? A: Materialized views are precomputed result sets stored in the database, while regular views

are virtual tables derived from base tables. Materialized views offer faster query performance but require periodic updates.

84. Q: What is the role of a data architect in database management? A: A data architect designs and plans the structure of databases, ensuring data integrity, scalability, and performance to meet business needs.
85. Q: What is the purpose of the ORDER BY clause in SQL? A: The ORDER BY clause is used to sort the result of a SQL query based on specified columns, in either ascending or descending order.
86. Q: What is a multi-table join in SQL? A: A multi-table join involves combining data from multiple tables using various join operations such as INNER JOIN, LEFT JOIN, RIGHT JOIN, and FULL OUTER JOIN.
87. Q: How does a database implement data replication across multiple servers? A: Databases use replication techniques like master-slave and master-master replication to copy data across multiple servers in real-time or near real-time.
88. Q: What is an entity in the context of the Entity-Relationship (ER) model? A: An entity represents a real-world object or concept that is stored as a table in a database, having attributes (columns) that describe its properties.
89. Q: What is the purpose of the COMMIT and ROLLBACK statements in transactions? A: The COMMIT statement is used to save the changes made during a transaction, while the ROLLBACK statement is used to undo the changes and restore the database to its state before the transaction.
90. Q: What is a data lake, and how does it differ from a data warehouse? A: A data lake is a storage repository that holds a vast amount of raw and unstructured data, while a data warehouse is designed to store structured and processed data for business intelligence and reporting purposes.
91. Q: What are NoSQL database categories, and what types of data do they handle? A: NoSQL databases are categorized into key-value stores, document stores, column-family stores, and graph databases, each designed to handle specific types of data and use cases.
92. Q: What is data governance in the context of database management? A: Data governance is a set of policies, procedures, and controls that ensure data quality, security, and compliance within an organization.
93. Q: What is the purpose of the TRUNCATE statement in SQL? A: The TRUNCATE statement is used to delete all records from a table quickly, but it cannot be rolled back like a DELETE statement.
94. Q: What is the role of an index in query optimization? A: An index improves query performance by reducing the number of data pages that need to be scanned during data retrieval.
95. Q: What is the purpose of the UNIQUE constraint in a database? A: The UNIQUE constraint ensures that the values in a column or a combination of columns are unique across all rows in the table.
96. Q: What is the difference between a database administrator (DBA) and a database developer? A: A DBA is responsible for managing the database system, while a database developer is involved in designing and implementing database applications and queries.
97. Q: How does a database handle concurrency control to prevent data inconsistencies? A: Databases use locking mechanisms, isolation levels, and transaction management to handle concurrent access and maintain data consistency.

98. Q: What is the role of a distributed database in managing data across multiple locations? A: A distributed database allows data to be stored and accessed across multiple physical locations, enabling better data availability and disaster recovery.
99. Q: What is a NoSQL graph database, and how is it useful? A: A NoSQL graph database uses graph structures to represent and store data, making it ideal for managing complex relationships and interconnected data.
100. Q: What is database replication lag, and how can it impact data consistency? A: Database replication lag refers to the delay in replicating data from the primary to the secondary server. High replication lag can lead to data inconsistencies between the servers.

Of course! Let's continue with more questions on database management systems:

101. Q: What is the purpose of the SQL UPDATE statement, and how is it used? A: The SQL UPDATE statement is used to modify existing records in a database table, allowing you to change the values of specific columns based on specified conditions.
102. Q: What is a NoSQL document store database, and when is it suitable for use? A: A NoSQL document store database stores data as documents, typically in JSON or BSON format. It is suitable for storing and querying semi-structured or unstructured data with varying attributes.
103. Q: What is database partitioning, and how does it improve performance? A: Database partitioning involves dividing a large table into smaller, more manageable partitions to enhance query performance, data maintenance, and storage management.
104. Q: How does the SQL UNION operator work, and when should you use it? A: The SQL UNION operator combines the result sets of two or more SELECT queries into a single result set. Use UNION when you want to merge data from multiple queries with the same structure.
105. Q: What is a database trigger, and what are its different types? A: A database trigger is a special type of stored procedure that automatically executes in response to specific database events such as INSERT, UPDATE, or DELETE. Types include BEFORE and AFTER triggers.
106. Q: What is the purpose of a database connection string? A: A database connection string contains the information required to establish a connection to a database, including the server address, username, password, and database name.
107. Q: What is data migration, and what are the challenges involved? A: Data migration is the process of transferring data from one database system to another. Challenges include data mapping, data cleansing, and ensuring data integrity during the migration.
108. Q: What is a NoSQL column-family store database, and how is it different from other NoSQL databases? A: A NoSQL column-family store database stores data in column families, where each row can have varying numbers of columns. It is suitable for large-scale data with varying attributes.
109. Q: How does database replication contribute to high availability and fault tolerance? A: Database replication creates copies of data on multiple servers, allowing for automatic failover and ensuring data availability even if one server experiences a failure.
110. Q: What is database caching, and how does it improve performance? A: Database caching stores frequently accessed data in memory, reducing the need to fetch data from disk, thereby improving query response times.

111. Q: What is database denormalization, and when should it be used? A: Database denormalization involves introducing redundancy into a database to improve performance by reducing the need for complex joins. It should be used cautiously to maintain data consistency.
112. Q: How does the SQL INSERT statement work, and what precautions should be taken while inserting data? A: The SQL INSERT statement adds new records to a table. Precautions include validating data before insertion, handling primary key constraints, and using parameterized queries to prevent SQL injection.
113. Q: What are database constraints, and how do they ensure data integrity? A: Database constraints impose rules on the data stored in a table to prevent invalid or inconsistent data from being inserted or updated, thus maintaining data integrity.
114. Q: What is a database index seek, and how does it differ from a database index scan? A: A database index seek uses the index to find specific data efficiently, while an index scan reads through the entire index to locate the required data.
115. Q: What is the purpose of the SQL DELETE statement, and what are the implications of using it? A: The SQL DELETE statement removes rows from a database table. Using it without a WHERE clause can lead to accidental deletion of all records in the table.
116. Q: How do database transactions help maintain data consistency? A: Database transactions ensure that a series of database operations are treated as a single unit of work, either all succeeding or all failing, to maintain data consistency.
117. Q: What is a NoSQL key-value store database, and when is it appropriate to use? A: A NoSQL key-value store database stores data as key-value pairs and is suitable for high-speed data access, caching, and storing unstructured or semi-structured data.
118. Q: How does a database handle data indexing for large datasets? A: Databases use various data structures like B-trees and hash tables for indexing, which efficiently organize and store index data for quick data retrieval.
119. Q: What is a NoSQL graph database, and when is it beneficial? A: A NoSQL graph database is designed to store and process data with complex relationships and interconnectedness. It is beneficial for applications involving social networks, recommendation systems, and network analysis.
120. Q: What is a database query plan, and how does it impact query performance? A: A database query plan is an execution plan created by the database optimizer to retrieve data efficiently. A well-optimized query plan can significantly improve query performance.
121. Q: What is the purpose of the SQL MERGE statement, and how is it used? A: The SQL MERGE statement is used to perform an "upsert" operation, which combines the functionality of both INSERT and UPDATE statements based on certain conditions.
122. Q: What are the different types of NoSQL data models, and how do they store data? A: The main types of NoSQL data models are document stores, key-value stores, column-family stores, and graph databases, each designed to store and retrieve data in unique ways.
123. Q: What is database materialization, and when is it beneficial? A: Database materialization involves creating and storing the results of complex queries or views as actual tables. It is beneficial for frequently used queries to reduce computation time.

124. Q: What is data profiling in database management? A: Data profiling is the process of analyzing and evaluating the quality of data in a database, identifying patterns, anomalies, and potential data issues.
125. Q: How does the SQL WHERE clause work, and what are its common operators? A: The SQL WHERE clause filters data based on specified conditions, and common operators include "=", "<>", "<", ">", "<=", ">=", "LIKE," and "IN."
126. Q: What is data modeling, and why is it essential for designing a database? A: Data modeling is the process of defining the structure and relationships of data in a database, which is essential for accurately representing business requirements and ensuring data integrity.
127. Q: What is a NoSQL key-range store database, and when is it suitable for use? A: A NoSQL key-range store database stores data as ordered key-value pairs, allowing efficient range queries over keys. It is suitable for applications involving time-series data and range-based queries.
128. Q: How does horizontal partitioning differ from vertical partitioning in a database? A: Horizontal partitioning divides a table's rows into smaller subsets, while vertical partitioning separates columns into different physical storage locations.
129. Q: What is database replication conflict resolution, and how is it handled? A: Database replication conflict resolution involves resolving conflicts that may arise when the same data is updated concurrently on multiple servers. Techniques like timestamps or conflict detection algorithms are used to resolve conflicts.
130. Q: What is the purpose of a database trigger in the context of auditing? A: Database triggers can be used for auditing changes to the database, allowing you to track modifications made to specific tables for security and compliance purposes.
131. Q: How does database clustering enhance performance and reliability? A: Database clustering involves grouping multiple database servers together to distribute data and processing tasks, which improves performance and provides fault tolerance.
132. Q: What is a NoSQL time-series database, and when is it used? A: A NoSQL time-series database is designed to efficiently store and analyze time-stamped data, making it suitable for applications like IoT sensor data, financial data, and logs.
133. Q: What is the purpose of the SQL SELECT INTO statement? A: The SQL SELECT INTO statement creates a new table and inserts the result set of a SELECT query into that new table.
134. Q: What is data warehousing schema design, and what are the commonly used schemas? A: Data warehousing schema design involves organizing data for efficient querying and reporting. Commonly used schemas include star schema, snowflake schema, and galaxy schema.
135. Q: What is a NoSQL columnar database, and when is it appropriate to use? A: A NoSQL columnar database stores data in columns rather than rows, making it suitable for analytics and scenarios where data is read and aggregated in column-based operations.
136. Q: How does the SQL LIKE operator work, and when is it used in queries? A: The SQL LIKE operator is used for pattern matching in queries, allowing you to search for data using partial or wildcard values.
137. Q: What is the purpose of a database transaction log, and how is it used in recovery? A: The database transaction log records all changes to the database, providing a point-in-time recovery mechanism to restore the database to a consistent state after a failure.

138. Q: What is NoSQL database sharding, and how does it help scale databases? A: NoSQL database sharding involves partitioning data across multiple servers to distribute the load and enable horizontal scaling for large datasets and high-performance applications.
139. Q: What are the best practices for securing a database system? A: Best practices include using strong passwords, limiting user privileges, encrypting sensitive data, keeping the database software up-to-date, and implementing regular backups.
140. Q: What is the role of a data warehouse architect in database management? A: A data warehouse architect designs and implements data warehouses to support business intelligence and analytics requirements, ensuring data accuracy and performance.
141. Q: What is a NoSQL wide-column store database, and when is it appropriate to use? A: A NoSQL wide-column store database stores data in column families, where each row can have a varying number of columns. It is suitable for handling large amounts of sparsely populated data.
142. Q: What is database query optimization, and how can it be achieved? A: Database query optimization involves improving the performance of queries through various techniques such as creating indexes, rewriting queries, and using caching.
143. Q: What is a NoSQL distributed database, and what challenges does it address? A: A NoSQL distributed database spreads data across multiple nodes, addressing challenges like scalability, fault tolerance, and high availability in large-scale applications.
144. Q: What is database encryption, and why is it important for data security? A: Database encryption is the process of converting sensitive data into ciphertext to protect it from unauthorized access. It ensures data confidentiality even if the database is compromised.
145. Q: How does the SQL GROUP BY clause work, and when is it used? A: The SQL GROUP BY clause groups rows based on specified columns and is typically used with aggregate functions like SUM, AVG, COUNT, etc.
146. Q: What is a NoSQL document store database, and when is it not recommended? A: A NoSQL document store database is suitable for flexible and unstructured data, but it may not be recommended for applications requiring complex querying or analytics.
147. Q: What is database data replication consistency, and how is it maintained? A: Database data replication consistency ensures that data is the same across all replicated servers by using techniques like two-phase commit or eventual consistency.
148. Q: What is a database view materialization, and how does it impact performance? A: Database view materialization involves creating a physical table from a view's definition. It can improve performance for complex views but may lead to increased storage requirements.
149. Q: What is the purpose of the SQL TRIGGER statement, and how is it used? A: The SQL TRIGGER statement is used to define a trigger—a special type of stored procedure that is automatically executed in response to specified database events.
150. Q: What is NoSQL tunable consistency, and how does it address performance and data availability trade-offs? A: NoSQL tunable consistency allows applications to configure the level of consistency they need based on performance and data availability requirements.

151. Q: What is a database connection pool, and how does it enhance performance? A: A database connection pool is a cache of established database connections that can be reused to reduce the overhead of creating new connections and improve performance.
152. Q: What is a NoSQL event store database, and when is it used? A: A NoSQL event store database is designed to store and manage event data, making it suitable for applications requiring event sourcing and event-driven architectures.
153. Q: What is the purpose of the SQL JOIN operation, and what are its different types? A: The SQL JOIN operation combines rows from multiple tables based on related columns. Different types include INNER JOIN, LEFT JOIN, RIGHT JOIN, and FULL OUTER JOIN.
154. Q: What is database auditing, and why is it essential for data governance? A: Database auditing involves tracking and recording changes to data and database activities for security, compliance, and data governance purposes.
155. Q: What is NoSQL eventual consistency, and how is it different from strong consistency? A: NoSQL eventual consistency allows data changes to propagate across nodes over time, while strong consistency ensures that data is immediately consistent across all nodes.
156. Q: What is database resource optimization, and how is it achieved? A: Database resource optimization involves efficiently using hardware resources like CPU, memory, and disk to ensure optimal database performance and responsiveness.
157. Q: What is a NoSQL in-memory database, and when is it suitable for use? A: A NoSQL in-memory database stores data primarily in RAM, offering extremely fast data access and is suitable for high-performance applications that require low-latency access.
158. Q: What is database change data capture, and how is it used? A: Database change data capture tracks and captures data changes in real-time, enabling applications to react to data updates and propagate changes to other systems.
159. Q: What is a database backup strategy, and what are the common backup methods? A: A database backup strategy outlines the plan for creating and storing database backups. Common methods include full backups, differential backups, and incremental backups.
160. Q: What is NoSQL tunable consistency, and how does it address performance and data availability trade-offs? A: NoSQL tunable consistency allows applications to configure the level of consistency they need based on performance and data availability requirements.
161. Q: What is a NoSQL time-to-live (TTL) feature, and how is it used? A: A NoSQL time-to-live (TTL) feature allows data to automatically expire after a specified period, which is useful for managing data retention and data cleanup in certain scenarios.
162. Q: What is database data archiving, and why is it important for long-term data storage? A: Database data archiving involves moving older or less frequently accessed data to secondary storage to free up space in the main database while still retaining access to historical records.
163. Q: What is NoSQL partitioning, and how does it contribute to scalability? A: NoSQL partitioning involves splitting data across multiple nodes, which allows the database to handle large datasets and handle read and write operations in parallel, contributing to scalability.

164. Q: What is database data masking, and how is it used to protect sensitive data? A: Database data masking is a technique that replaces sensitive data with fictional or anonymized data to protect sensitive information from unauthorized access.
165. Q: What is the purpose of the SQL SET statement, and how is it used? A: The SQL SET statement is used to assign values to variables that can be used in SQL queries, making queries more dynamic and flexible.
166. Q: What is a NoSQL data consistency model, and what are the different consistency levels? A: A NoSQL data consistency model defines the guarantees provided by the database in terms of data consistency across replicas. Different levels include eventual consistency, strong consistency, and more.
167. Q: What is database connection pooling, and how does it benefit performance and resource utilization? A: Database connection pooling involves reusing existing database connections instead of creating new ones, which reduces connection overhead and enhances performance and resource efficiency.
168. Q: What is NoSQL read and write quorum, and how are they used in distributed databases? A: NoSQL read and write quorums determine the number of replicas that must acknowledge read or write operations for data consistency and availability in distributed databases.
169. Q: What is database high availability, and how is it achieved in a database system? A: Database high availability refers to the ability of a database system to remain operational and accessible even in the face of hardware failures, software issues, or network disruptions. It is achieved through redundancy, failover mechanisms, and replication.
170. Q: What is the purpose of the SQL DELETE CASCADE option, and when is it useful? A: The SQL DELETE CASCADE option automatically deletes related records in child tables when a record in the parent table is deleted. It is useful to maintain referential integrity and handle cascading deletes.
171. Q: What is NoSQL ACID compliance, and how is it different from ACID in SQL databases? A: NoSQL ACID compliance refers to the support for Atomicity, Consistency, Isolation, and Durability in NoSQL databases. However, the implementation of ACID properties in NoSQL databases may vary from traditional SQL databases.
172. Q: What is database partitioning key selection, and how is it determined? A: Database partitioning key selection involves choosing the column or attribute used to determine how data is distributed across partitions. It is typically based on the access patterns and the nature of the data.
173. Q: What is NoSQL eventual consistency, and how is it implemented in distributed databases? A: NoSQL eventual consistency allows data to be eventually consistent across replicas after a period of time. It is achieved through asynchronous replication and conflict resolution mechanisms.
174. Q: What is database hot standby, and how does it contribute to database availability? A: Database hot standby involves maintaining a synchronized replica of the primary database that can take over in case of a failure, ensuring minimal downtime and high availability.
175. Q: What is the purpose of the SQL ALTER TABLE statement, and how is it used? A: The SQL ALTER TABLE statement is used to modify the structure of an existing table, such as adding or dropping columns, changing data types, or adding constraints.

176. Q: What is NoSQL consistency tuning, and how does it impact performance and data integrity? A: NoSQL consistency tuning allows applications to configure the level of data consistency needed based on performance requirements, balancing consistency with potential data staleness.
177. Q: What is database data deduplication, and why is it beneficial for storage optimization? A: Database data deduplication is the process of identifying and eliminating duplicate data in a database, which reduces storage requirements and improves data efficiency.
178. Q: What is NoSQL sharding key selection, and how is it determined? A: NoSQL sharding key selection involves choosing the attribute or column used to determine how data is partitioned across shards in a distributed database. It should be selected based on even data distribution and access patterns.
179. Q: What is database horizontal scaling, and how is it achieved in distributed databases? A: Database horizontal scaling involves adding more servers or nodes to a distributed database system to handle increased data volume and workload. It can be achieved through partitioning and replication.
180. Q: What is the purpose of the SQL CASE statement, and how is it used? A: The SQL CASE statement allows conditional logic in SQL queries, enabling you to perform different actions based on specified conditions.
181. Q: What is NoSQL query optimization, and how does it differ from SQL query optimization? A: NoSQL query optimization involves optimizing queries for distributed databases, where data may be spread across multiple nodes. It differs from SQL query optimization, which is focused on optimizing queries for a single, centralized database.
182. Q: What is database table partitioning, and how does it improve data retrieval performance? A: Database table partitioning involves dividing a large table into smaller, more manageable partitions. It improves data retrieval performance by reducing the amount of data that needs to be scanned for specific queries.
183. Q: What is the CAP theorem, and how does it impact distributed database design? A: The CAP theorem states that a distributed database system cannot simultaneously guarantee all three properties of Consistency, Availability, and Partition Tolerance. Database designers must choose a trade-off between these properties based on their specific requirements.
184. Q: What is database data compression, and how is it applied to save storage space? A: Database data compression is the process of reducing the size of data stored in a database to save storage space. It is achieved by using algorithms that eliminate redundancies and use more efficient data representations.
185. Q: What is ACID-BASE, and how does it relate to database consistency in NoSQL systems? A: ACID-BASE is an alternative model to the traditional ACID model for databases. It relaxes some of the ACID properties to achieve better scalability and availability in NoSQL databases, accepting eventual consistency instead of immediate consistency.
186. Q: What is the purpose of the SQL HAVING clause, and when is it used in queries? A: The SQL HAVING clause filters data based on specified conditions after a GROUP BY operation, typically used with aggregate functions in GROUP BY queries.

187. Q: What is database data replication topology, and how does it affect replication performance and reliability? A: Database data replication topology refers to the way data is copied and distributed across replica nodes. Different topologies (e.g., master-slave, master-master) impact replication performance, latency, and data consistency.
188. Q: What is NoSQL query languages, and how do they differ from SQL? A: NoSQL databases often use query languages specific to their data model, which can differ significantly from SQL in terms of syntax and capabilities.
189. Q: What is database data tiering, and how is it utilized for cost optimization? A: Database data tiering involves categorizing data based on its access frequency and moving it to different storage tiers (e.g., fast SSD, slower HDD) to optimize costs and performance.
190. Q: What is NoSQL serverless architecture, and how does it benefit database management? A: NoSQL serverless architecture allows developers to focus on application logic without managing the underlying infrastructure, providing automatic scaling and cost optimization.
191. Q: What is database in-memory processing, and how does it improve query performance? A: Database in-memory processing stores data in RAM for faster data access and processing, significantly improving query performance compared to disk-based databases.
192. Q: What is the role of a database performance tuning expert, and how do they optimize database performance? A: A database performance tuning expert analyzes database configurations, queries, indexes, and hardware resources to identify and implement optimizations that improve overall database performance.
193. Q: What is NoSQL geospatial database, and when is it suitable for use? A: A NoSQL geospatial database is designed to handle spatial data, enabling queries based on geographical locations, making it suitable for mapping, GIS, and location-based applications.
194. Q: What is database buffer pool management, and how does it impact data caching? A: Database buffer pool management involves managing the cache of data pages in memory to improve query performance by reducing the need to read data from disk.
195. Q: What is NoSQL change data capture, and how is it used for real-time data processing? A: NoSQL change data capture captures and streams data changes in real-time, allowing applications to react to data updates and propagate changes to downstream systems.
196. Q: What is database table indexing strategy, and how is it determined? A: Database table indexing strategy involves choosing the right columns and types of indexes to create for a table, based on the types of queries and access patterns performed on the data.
197. Q: What is NoSQL database compaction, and how does it benefit data storage efficiency? A: NoSQL database compaction is the process of reclaiming storage space by removing obsolete or deleted data and optimizing data organization, resulting in improved storage efficiency.
198. Q: What is database read scaling, and how is it achieved in distributed databases? A: Database read scaling involves distributing read queries across multiple replica nodes to handle read-intensive workloads and increase query throughput in a distributed database.
199. Q: What is NoSQL BigTable database, and what are its use cases? A: NoSQL BigTable database is a distributed, wide-column store database designed for large-scale,

high-performance applications, such as Google's services, analytics platforms, and IoT data processing.

200. Q: What is the purpose of the SQL CROSS JOIN statement, and when is it used in queries? A: The SQL CROSS JOIN statement produces the Cartesian product of two or more tables, combining all rows from each table. It is used when you need to combine all possible combinations of rows between tables.
201. Q: What is Oracle Database? A: Oracle Database is a relational database management system (RDBMS) developed by Oracle Corporation. It is widely used for storing, managing, and retrieving data in various applications.
202. Q: What are the main components of an Oracle database? A: The main components are the instance (memory structures and background processes) and the database (data files, control files, and redo log files).
203. Q: What is an Oracle tablespace? A: An Oracle tablespace is a logical storage container that groups related data files, providing better management of storage space and performance.
204. Q: Explain the difference between a unique key and a primary key in Oracle. A: A unique key ensures that the values in a column are unique, but it allows for NULL values. A primary key is a unique key that also enforces a NOT NULL constraint.
205. Q: How can you create a new user in Oracle? A: You can create a new user using the CREATE USER statement and specifying the username and password, along with other optional settings like default tablespace and temporary tablespace.
206. Q: What is the purpose of the Oracle Data Dictionary? A: The Oracle Data Dictionary is a collection of database tables and views that store metadata about the database, including information about tables, columns, indexes, users, and privileges.
207. Q: Explain the difference between a clustered and non-clustered index in Oracle. A: A clustered index determines the physical order of data in a table, whereas a non-clustered index creates a separate structure to map the index to the data.
208. Q: What is an Oracle sequence, and how is it used? A: An Oracle sequence is a database object that generates unique integers in a sequential order. It is commonly used to generate primary key values for tables.
209. Q: How can you backup an Oracle database, and what are the different backup methods? A: Oracle databases can be backed up using the RMAN (Recovery Manager) utility. Different backup methods include full backups, incremental backups, and archive log backups.
210. Q: What is the purpose of the Oracle RAC (Real Application Clusters)? A: Oracle RAC allows multiple instances to access a single Oracle database simultaneously, providing high availability, scalability, and load balancing.
211. Q: How can you tune the performance of an Oracle database? A: Performance tuning in Oracle involves optimizing SQL queries, creating indexes, adjusting memory settings, and analyzing execution plans.

212. Q: What is an Oracle PL/SQL block, and how is it executed? A: An Oracle PL/SQL block is a set of one or more SQL and PL/SQL statements that are executed together as a unit. It can be executed using anonymous blocks or stored procedures/functions.
213. Q: How can you find the number of rows affected by an SQL statement in Oracle? A: You can use the SQL%ROWCOUNT attribute to retrieve the number of rows affected by the last SQL statement.
214. Q: Explain the difference between COMMIT and ROLLBACK in Oracle. A: COMMIT is used to save the changes made during a transaction, while ROLLBACK is used to undo the changes and restore the database to its state before the transaction.
215. Q: What is the purpose of the Oracle *SQLPlus* tool? A: *SQLPlus* is a command-line interface tool provided by Oracle for interacting with the Oracle Database. It is widely used for executing SQL queries and scripts.
216. Q: What is an Oracle trigger, and when is it used? A: An Oracle trigger is a named PL/SQL block that is automatically executed when specific database events occur, such as INSERT, UPDATE, or DELETE on a table.
217. Q: How can you find the execution plan for an SQL statement in Oracle? A: You can use the EXPLAIN PLAN statement or the DBMS_XPLAN package to view the execution plan for an SQL statement.
218. Q: What are Oracle hints, and how are they used in SQL queries? A: Oracle hints are special comments that provide instructions to the Oracle optimizer on how to execute a specific SQL statement. They are used to influence the execution plan.
219. Q: Explain the difference between a VARCHAR and a CHAR data type in Oracle. A: Both VARCHAR and CHAR are used to store character data. The main difference is that VARCHAR stores variable-length strings, while CHAR stores fixed-length strings.
220. Q: How can you grant privileges to a user in Oracle? A: You can use the GRANT statement to grant privileges such as SELECT, INSERT, UPDATE, DELETE, and others to a user on specific tables or database objects.

1) Define Database.

A prearranged collection of figures known as data is called database.

2) What is DBMS?

Database Management Systems (DBMS) are applications designed especially which enable user interaction with other applications.

Or, The database management system is a collection of programs that enables user to store, retrieve, update and delete information from a database.

3) What are the various kinds of interactions catered by DBMS?

The various kind of interactions catered by DBMS are:

- Data definition
- Update
- Retrieval
- Administration

4) What is RDBMS ?

Relational Database Management system (RDBMS) is a database management system (DBMS) that is based on the relational model. Data from relational database can be accessed or reassembled in many different ways without having to reorganize the database tables. Data from relational database can be accessed using an API, Structured Query Language (SQL).

5) Segregate database technology's development.

The development of database technology is divided into:

- Structure or data model
- Navigational model
- SQL/ relational model

6) Who proposed the relational model?

Edgar F. Codd proposed the relational model in 1970.

7) What are the features of Database language?

A database language may also incorporate features like:

DBMS-specific Configuration and management of storage engine

Computations to modification of query results by computations, like summing, counting,

averaging, grouping, sorting and cross-referencing Constraint enforcement Application Programming Interface

8) What do database languages do?

As special-purpose languages, they have:

- Data definition language
- Data manipulation language
- Query language

9) Define database model.

- A data model determining fundamentally how data can be stored, manipulated and organised and the structure of the database logically is called database model.

10) What is SQL?

Structured Query Language (SQL) being ANSI standard language updates database and commands for accessing.

Or, Structured Query Language(SQL) is a language designed specifically for communicating with databases. SQL is an ANSI (American National Standards Institute) standard.

11) What are the different type of SQL's statements ?

This is one of the most frequently asked SQL Interview Questions for freshers. SQL statements are broadly classified into three. They are

1. DDL – Data Definition Language

DDL is used to define the structure that holds the data. For example, Create, Alter, Drop and Truncate table.

2. DML – Data Manipulation Language

DML is used for manipulation of the data itself. Typical operations are Insert, Delete, Update and retrieving the data from the table. The Select statement is considered as a limited version of the DML, since it can't change the data in the database. But it can perform operations on data retrieved from the DBMS, before the results are returned to the calling function.

3. DCL – Data Control Language

DCL is used to control the visibility of data like granting database access and set privileges to create tables, etc. Example - Grant, Revoke access permission to the user to access data in the database.

12) Enlist some commands of DDL.

They are:

CREATE:

Create is used in the CREATE TABLE statement. Syntax is:

CREATE TABLE [column name] ([column definitions]) [table parameters]

ALTER:

It helps in modification of an existing object of database. Its syntax is:

ALTER objecttype objectname parameters.

DROP:

It destroys an existing database, index, table or view. Its syntax is:

DROP objecttype objectname.

13) What are the Advantages of SQL ?

1. **SQL is not a proprietary language** used by specific database vendors. Almost every major DBMS supports SQL, so learning this one language will enable programmers to interact with any database like ORACLE, SQL ,MYSQL etc.
2. **SQL is easy to learn.** The statements are all made up of descriptive English words, and there aren't that many of them.
3. SQL is actually a very powerful language and by using its language elements you can perform very **complex and sophisticated database operations**.

14) What is a field in a database ?

A field is an area within a record reserved for a specific piece of data.

Examples: Employee Name, Employee ID, etc.

15) Enlist the various relationships of database.

The various relationships of database are:

- One-to-one: Single table having drawn relationship with another table having similar kind of columns.
- One-to-many: Two tables having primary and foreign key relation.
- Many-to-many: Junction table having many tables related to many tables.

16) What is a database transaction?

Database transaction takes database from one consistent state to another. At the end of the transaction the system must be in the prior state if the transaction fails or the status of the system should reflect the successful completion if the transaction goes through.

17) What are properties of a transaction?

Expect this SQL Interview Questions as a part of an any interview, irrespective of your experience. Properties of the transaction can be summarized as ACID Properties.

1. Atomicity

A transaction consists of many steps. When all the steps in a transaction get completed, it will get reflected in DB or if any step fails, all the transactions are rolled back.

2. Consistency

The database will move from one consistent state to another, if the transaction succeeds and remain in the original state, if the transaction fails.

3. Isolation

Every transaction should operate as if it is the only transaction in the system.

4. Durability

Once a transaction has completed successfully, the updated rows/records must be available for all other transactions on a permanent basis.

18) What is a Database Lock ?

Database lock tells a transaction, if the data item in questions is currently being used by other transactions.

19) Define Normalization.

Organized data void of inconsistent dependency and redundancy within a database is called normalization.

20) What are the different type of normalization?

In database design, we start with one single table, with all possible columns. A lot of redundant data would be present since it's a single table. The process of removing the redundant data, by splitting up the table in a well defined fashion is called normalization.

1. First Normal Form (1NF)

A relation is said to be in first normal form if and only if all underlying domains contain atomic values only. After 1NF, we can still have redundant data.

2. Second Normal Form (2NF)

A relation is said to be in 2NF if and only if it is in 1NF and every non key attribute is fully dependent on the primary key. After 2NF, we can still have redundant data.

3. Third Normal Form (3NF)

A relation is said to be in 3NF, if and only if it is in 2NF and every non key attribute is non-transitively dependent on the primary key.

21) Enlist the advantages of normalizing database.

Advantages of normalizing database are:

- No duplicate entries
- Saves storage space
- Boasts the query performances.

22) Define Denormalization.

Boosting up database performance, adding of redundant data which in turn helps rid of complex data is called denormalization.

23) What is a primary key?

A primary key is a column whose values **uniquely identify every row** in a table. Primary key values can never be reused. If a row is deleted from the table, its primary key may not be assigned to any new rows in the future. To define a field as primary key, following conditions had to be met :

1. No two rows can have the same primary key value.
2. Every row must have a primary key value.
3. The primary key field cannot be null.
4. Value in a primary key column can never be modified or updated, if any foreign key refers to that primary key.

24) What is a Composite Key ?

A Composite primary key is a type of candidate key, which represents a set of columns whose values uniquely identify every row in a table.

For example - if "Employee_ID" and "Employee Name" in a table is combined to uniquely identify a row its called a Composite Key.

25) What is a Foreign Key ?

When a "one" table's primary key field is added to a related "many" table in order to create the common field which relates the two tables, it is called a foreign key in the "many" table.

For example, the salary of an employee is stored in salary table. The relation is established via foreign key column "Employee_ID_Ref" which refers "Employee_ID" field in the Employee table.

26) What is a Unique Key ?

Unique key is same as primary with the difference being the existence of null. Unique key field allows one value as NULL value.

27) Define Union All operator and Union.

Full recordings of two tables is Union All operator.

A distinct recording of two tables is Union.

28) Define cursor.

A database object which helps in manipulating data row by row representing a result set is called cursor.

29) Enlist the cursor types.

They are:

- Dynamic: it reflects changes while scrolling.
- Static: doesn't reflect changes while scrolling and works on recording of snapshot.
- Keyset: data modification without reflection of new data is seen.

30) Enlist the types of cursor.

They types of cursor are:

- Implicit cursor: Declared automatically as soon as the execution of SQL takes place without the awareness of the user.
- Explicit cursor: Defined by PL/ SQL which handles query in more than one row.

31) Define sub-query.

A query contained by a query is called Sub-query.

32) Why is group-clause used?

Group-clause uses aggregate values to be derived by collecting similar data.

33) Compare Non-clustered and clustered index

Both having B-tree structure, non-clustered index has data pointers enabling one table many non-clustered indexes while clustered index is distinct for every table.

34) Define Aggregate functions.

Functions which operate against a collection of values and returning single value is called aggregate functions

35) Define Scalar functions.

Scalar function is depended on the argument given and returns sole value.

36) What is a view?

The views are virtual tables. Unlike tables that contain data, views simply contain queries that dynamically retrieve data when used.

37) What is a materialized view?

Materialized views are also a view but are disk based. **Materialized views** get updates on specific duration, base upon the interval specified in the query definition. We can index materialized view.

38) What are the advantages and disadvantages of views in a database?

Advantages:

1. Views don't store data in a physical location.
2. The view can be used to hide some of the columns from the table.
3. Views can provide Access Restriction, since data insertion, update and deletion is not possible with the view.

Disadvantages:

1. When a table is dropped, associated view become irrelevant.
2. Since the view is created when a query requesting data from view is triggered, its a bit slow.
3. When views are created for large tables, it occupies more memory.

39) What restrictions can you apply when you are creating views?

Restrictions that are applied are:

- Only the current database can have views.
- You are not liable to change any computed value in any particular view.
- Integrity constants decide the functionality of INSERT and DELETE.
- Full-text index definitions cannot be applied.
- Temporary views cannot be created.
- Temporary tables cannot contain views.
- No association with DEFAULT definitions.
- Triggers such as INSTEAD OF is associated with views.

40) Define “correlated subqueries”.

A ‘correlated subquery’ is a sort of sub query but correlated subquery is reliant on another query for a value that is returned. In case of execution, the sub query is executed first and then the correlated query.

41) Define Data Warehousing.

Storage and access of data from the central location in order to take some strategic decision is called Data Warehousing. Enterprise management is used for managing the information whose framework is known as Data Warehousing.

42) Define Join and enlist its types.

Joins help in explaining the relation between different tables. They also enable you to select data with relation to data in another table.

The various types are:

- INNER JOINS: Blank rows are left in the middle while more than equal to two tables are joined.
- OUTER JOINS: Divided into Left Outer Join and Right Outer Join. Blank rows are left at the specified side by joining tables in other side.

Other joins are CROSS JOINS, NATURAL JOINS, EQUI JOIN and NON-EQUI JOIN.

Cross Join will return all records where each row from the first table is combined with each row from the second table.

43) What do you mean by Index hunting?

Indexes help in improving the speed as well as the query performance of database. The procedure of boosting the collection of indexes is named as Index hunting.

44) How does Index hunting help in improving query performance?

Index hunting helps in improving the speed as well as the query performance of database. The followed measures are achieved to do that:

- The query optimizer is used to coordinate the study of queries with the workload and the best use of queries suggested based on this.
- Index, query distribution along with their performance is observed to check the effect.
- Tuning databases to a small collection of problem queries is also recommended.

45) Enlist the disadvantages of query.

The disadvantages of query are:

- No indexes

- Stored procedures are excessively compiled.
- Triggers and procedures are without SET NOCOUNT ON.
- Complicated joins making up inadequately written query.
- Cursors and temporary tables showcase a bad presentation.

46) Enlist ways to efficiently code transactions.

Ways to efficiently code transactions:

- User input should not be allowed while transactions.
- While browsing, transactions must not be opened of data.
- Transactions must be kept as small as possible.
- Lower transaction segregation levels.
- Least information of data must be accessed while transacting.

47) What is Executive Plan?

Executive plan can be defined as:

- SQL Server caches collected procedure or the plan of query execution and used thereafter by subsequent calls.
- An important feature in relation to performance enhancement.
- Data execution plan can be viewed textually or graphically.

48) Define B-trees.

A data structure in the form of tree which stores sorted data and searches, insertions, sequential access and deletions are allowed in logarithmic time.

49) Differentiate Table Scan from Index Scan.

Iterating over all the table rows is called Table Scan while iterating over all the index items is defined as Index Scan.

50) What do you mean by Fill Factor concept with respect to indexes?

Fill Factor can be defined as being that value which defines the percentage of left space on every leaf-level page that is to be packed with data. 100 is the default value of Fill Factor.

51) Define Fragmentation.

Fragmentation can be defined as a database feature of server that promotes control on data which is stored at table level by the user.

52) Differentiate Nested Loop, Hash Join and Merge Join.

Nested loop (loop over loop)

An outer loop within an inner loop is formed consisting of fewer entries and then for individual entry, inner loop is individually processed.

E.g.

- `Select col1.*, col2.* from coll, col2 where coll.col1=col2.col2;`

It's processing takes place in this way:

For i in (select * from col1) loop

For j in (select * from col2 where col2=i.col1) loop

Results are displayed;

End of the loop;

End of the loop;

The Steps of nested loop are:

- Identify outer (driving) table
- Assign inner (driven) table to outer table.
- For every row of outer table, access the rows of inner table.

Nested Loops is executed from the inner to the outer as:

- outer_loop
- inner_loop
- Hash join

While joining large tables, the use of Hash Join is preferred.

Algorithm of Hash Join is divided into:

- Build: It is a hash table having in-memory which is present on the smaller table.
- Probe: this hash value of the hash table is applicable for each second row element.
- Sort merge join

Two independent sources of data are joined in sort merge join. Their performance is better as compared to nested loop when the data volume is big enough but it is not good as hash joins generally.

The full operation can be divided into parts of two:

Sort join operation :

Get first row R1 from input1

Get first row R2 from input2.

Merge join operation:

‘while’ is not present at either loop’s end.

if R1 joins with R2

next row is got R2 from the input 2

return (R1, R2)

else if R1 < style=""> next row is got from R1 from input 1

else
next row is got from R2 from input 2
end of the loop

53) What is Database partitioning?

Division of logical database into independent complete units for improving its management, availability and performance is called Database partitioning.

54) Explain the importance of partitioning.

Splitting of one table which is large into smaller database entities logically is called database partitioning. Its benefits are:

- To improve query performance in situations dramatically when mostly rows which are heavily accessed are in one partition.
- Accessing large parts of a single partition
- Slower and cheaper storage media can be used for data which is seldom used.

55) Define Database system.

DBMS along with database is called Database system.

56) What do you mean by Query Evaluation Engine?

Query Evaluation Engine executes the low-level instructions that are generated by the compiler.

57) Define DDL Interpreter.

DDL statements are interpreted and recorded in tables called metadata.

58) Define Atomicity and Aggregation.

Atomicity: It's an all or none concept which enables the user to be assured of incomplete transactions to be taken care of. The actions involving incomplete transactions are left undone in DBMS.

Aggregation: The collected entities and their relationship are aggregated in this model. It is mainly used in expressing relationships within relationships.

59) Enlist the various transaction phases.

The various transaction phases are:

- Analysis Phase.
- Redo Phase
- Undo Phase

60) Define Object-oriented model.

Compilations of objects make up this model in which values are stored within instance variables which is inside the object. The object itself comprises bodies of object for its operation which are called methods. Objects containing same kind of variables and methods are called classes.

61) Define Entity.

It can be defined as being a 'thing' with an independent existence in the real world.

62) What do you mean by Entity type?

A set of entries having similar attributes are entity types.

63) Define Entity Set.

Compilation of all entries of any particular type of entry in the database is called Entity Set.

64) What do you mean by Entity type extension?

Compilation of similar entity types into one particular type which is grouped together as an entity set.

65) What is a stored procedure?

Stored Procedure is a function which contains a collection of SQL Queries. The procedure can take inputs, process them and send back output.

66) What are the advantages of a stored procedure?

Stored Procedures are precompiled and stored in the database. This enables the database to execute the queries much faster. Since many queries can be included in a stored procedure, round trip time to execute multiple queries from source code to database and back is avoided.

67) What is a trigger?

Database triggers are sets of commands that get executed when an event (Before Insert, After Insert, On Update, On delete of a row) occurs on a table, views.

68) Explain the difference between DELETE, TRUNCATE and DROP commands?

Once **delete operation** is performed, Commit and Rollback can be performed to retrieve data.

Once the **truncate** statement is executed, Commit and Rollback statement cannot be performed. Where condition can be used along with delete statement but it can't be used with

truncate statement.

Drop command is used to drop the table or keys like primary,foreign from a table.

69) What is the difference between Cluster and Non cluster Index?

A **clustered index** reorders the way records in the table are physically stored. There can be only one clustered index per table. It makes data retrieval faster.

A **non clustered index** does not alter the way it was stored but creates a completely separate object within the table. As a result insert and update command will be faster

70) What is Union, minus and Intersect commands?

MINUS operator is used to return rows from the first query but not from the second query.
INTERSECT operator is used to return rows returned by both the queries

71) What is SQL Profiler?

Microsoft SQL Server Profiler is a graphical user interface to SQL Trace for monitoring an instance of the Database Engine or Analysis Services. You can capture and save data about each event to a file or table to analyze later.

Use SQL Profiler to monitor only the events in which you are interested.

If traces are becoming too large, you can filter them based on the information you want, so that only a subset of the event data is collected. Monitoring too many events adds overhead to the server and the monitoring process and can cause the trace file or trace table to grow very large, especially when the monitoring process takes place over a long period of time.

MYSQL Interview Questions & Answers

72) What is MySQL?

MySQL is an open source DBMS which is built, supported and distributed by MySQL AB (now acquired by Oracle)

73) What are the technical features of MySQL?

MySQL database software is a client or server system which includes

- Multithreaded SQL server supporting various client programs and libraries
- Different backend
- Wide range of application programming interfaces and
- Administrative tools.

74) Why MySQL is used?

MySQL database server is reliable, fast and very easy to use. This software can be downloaded as freeware and can be downloaded from the internet.

75) What are the advantages of MySQL when compared with Oracle?

MySQL is open source software which is available at any time and has no cost involved.

- MySQL is portable
- GUI with command prompt.
- Administration is supported using MySQL Query Browser

76) Differentiate between FLOAT and DOUBLE?

Following are differences for FLOAT and DOUBLE:

- Floating point numbers are stored in FLOAT with eight place accuracy and it has four bytes.
- Floating point numbers are stored in DOUBLE with accuracy of 18 places and it has eight bytes.

77) . Differentiate CHAR_LENGTH and LENGTH?

CHAR_LENGTH is character count whereas the LENGTH is byte count. The numbers are same for Latin characters but they are different for Unicode and other encodings.

78) How to represent ENUMs and SETs internally?

ENUMs and SETs are used to represent powers of two because of storage optimizations.

79) What is the usage of ENUMs in MySQL?

ENUM is a string object used to specify set of predefined values and that can be used during table creation.

Create table size(name ENUM('Small', 'Medium','Large'));

80) Define REGEXP?

REGEXP is a pattern match in which matches pattern anywhere in the search value.

81) Difference between CHAR and VARCHAR?

Following are the differences between CHAR and VARCHAR:

- CHAR and VARCHAR types differ in storage and retrieval
- CHAR column length is fixed to the length that is declared while creating table. The length value ranges from 1 and 255
- When CHAR values are stored then they are right padded using spaces to specific length. Trailing spaces are removed when CHAR values are retrieved.

82) Give string types available for column?

The string types are:

- SET

- BLOB
- ENUM
- CHAR
- TEXT
- VARCHAR

83) How to get current MySQL version?

SELECT VERSION (); is used to get the current version of MySQL.

84) What storage engines are used in MySQL?

Storage engines are called table types and data is stored in files using various techniques.

Technique involves:

- Storage mechanism
- Locking levels
- Indexing
- Capabilities and functions.

85) What are the drivers in MySQL?

Following are the drivers available in MySQL:

- PHP Driver
- JDBC Driver
- ODBC Driver
- C WRAPPER
- PYTHON Driver
- PERL Driver
- RUBY Driver
- CAP11PHP Driver
- Ado.net5.mxj

86) What is the difference between primary key and candidate key?

Every row of a table is identified uniquely by primary key. There is only one primary key for a table.

Primary Key is also a candidate key. By common convention, candidate key can be designated as primary and which can be used for any foreign key references.

87) What are the column comparisons operators?

The =, <>, <=, <, >=, >, <<, >>, <=>, AND, OR, or LIKE operators are used in column comparisons in SELECT statements.

88) How can we get the number of rows affected by query?

Number of rows can be obtained by SELECT COUNT (user_id) FROM users;

89) Is Mysql query is case sensitive?

NO.

90) What are the different tables present in MySQL?

Total 5 types of tables are present:

- MyISAM
- Heap
- Merge
- INNO DB
- ISAM

MyISAM is the default storage engine as of MySQL.

91) What is ISAM?

ISAM is abbreviated as Indexed Sequential Access Method. It was developed by IBM to store and retrieve data on secondary storage systems like tapes.

92) How to display top 50 rows?

In MySQL, top 50 rows are displayed by using this following query:

```
1 SELECT * FROM  
2 LIMIT 0,50;
```

93) How many TRIGGERS are allowed in MySQL table?

SIX triggers are allowed in MySQL table. They are as follows:

- BEFORE INSERT
- AFTER INSERT
- BEFORE UPDATE
- AFTER UPDATE
- BEFORE DELETE and
- AFTER DELETE

Oracle Interview Questions & Answers

94) Difference between varchar and varchar2 data types?

Varchar can store up to 2000 bytes and varchar2 can store up to 4000 bytes. Varchar will occupy space for NULL values and Varchar2 will not occupy any space. Both are differed with respect to space.

95) In which language Oracle has been developed?

Oracle has been developed using C Language.

96) What is RAW datatype?

RAW datatype is used to store values in binary data format. The maximum size for a raw in a table is 32767 bytes.

97) What is the use of NVL function?

The NVL function is used to replace NULL values with another or given value. Example is –

NVL(Value, replace value)

98) Whether any commands are used for Months calculation? If so, What are they?

In Oracle, months_between function is used to find number of months between the given dates. Example is –Months_between(Date 1, Date 2)

99) What are nested tables?

Nested table is a data type in Oracle which is used to support columns containing multi valued attributes. It also hold entire sub table.

100) What is COALESCE function?

COALESCE function is used to return the value which is set to be not null in the list. If all values in the list are null, then the coalesce function will return NULL.

Coalesce(value1, value2,value3,...)

101) What is BLOB datatype?

A BLOB data type is a varying length binary string which is used to store two gigabytes memory. Length should be specified in Bytes for BLOB.

102) How do we represent comments in Oracle?

Comments in Oracle can be represented in two ways –

1. Two dashes(–) before beginning of the line – Single statement
2. /*—— */ is used to represent it as comments for block of statement
- 3.

103) What is DML?

Data Manipulation Language (DML) is used to access and manipulate data in the existing objects. DML statements are insert, select, update and delete and it won't implicitly commit the current transaction.

104) What is the difference between TRANSLATE and REPLACE?

Translate is used for character by character substitution and Replace is used substitute a single character with a word.

105) How do we display rows from the table without duplicates?

Duplicate rows can be removed by using the keyword DISTINCT in the select statement.

106) What is the usage of Merge Statement?

Merge statement is used to select rows from one or more data source for updating and insertion into a table or a view. It is used to combine multiple operations.

107) What is NULL value in oracle?

NULL value represents missing or unknown data. This is used as a place holder or represented it in as default entry to indicate that there is no actual data present.

108) What is USING Clause and give example?

The USING clause is used to specify with the column to test for equality when two tables are joined.

```
[sql]Select * from employee join salary using employee ID[/sql]
```

Employee tables join with the Salary tables with the Employee ID.

109) What is key preserved table?

A table is set to be key preserved table if every key of the table can also be the key of the result of the join. It guarantees to return only one copy of each row from the base table.

110) What is WITH CHECK OPTION?

The WITH CHECK option clause specifies check level to be done in DML statements. It is used to prevent changes to a view that would produce results that are not included in the sub query.

111) What is the use of Aggregate functions in Oracle?

Aggregate function is a function where values of multiple rows or records are joined together to get a single value output. Common aggregate functions are –

- Average
- Count
- Sum

112) What do you mean by GROUP BY Clause?

A GROUP BY clause can be used in select statement where it will collect data across multiple records and group the results by one or more columns.

113) What is a sub query and what are the different types of subqueries?

Sub Query is also called as Nested Query or Inner Query which is used to get data from multiple tables. A sub query is added in the where clause of the main query.

There are two different types of subqueries:

- Correlated sub query

A Correlated sub query cannot be as independent query but can reference column in a table listed in the from list of the outer query.

- Non-Correlated subquery

This can be evaluated as if it were an independent query. Results of the sub query are submitted to the main query or parent query.

114) What is cross join?

Cross join is defined as the Cartesian product of records from the tables present in the join. Cross join will produce result which combines each row from the first table with the each row from the second table.

115) What are temporal data types in Oracle?

Oracle provides following temporal data types:

- Date Data Type – Different formats of Dates
- TimeStamp Data Type – Different formats of Time Stamp
- Interval Data Type – Interval between dates and time

116) How do we create privileges in Oracle?

A privilege is nothing but right to execute an SQL query or to access another user object. Privilege can be given as system privilege or user privilege.

```
[sql]GRANT user1 TO user2 WITH MANAGER OPTION;[/sql]
```

117) What is VArray?

VArray is an oracle data type used to have columns containing multivalued attributes and it can hold bounded array of values.

118) How do we get field details of a table?

Describe <Table_Name> is used to get the field details of a specified table.

119) What is the difference between rename and alias?

Rename is a permanent name given to a table or a column whereas Alias is a temporary name given to a table or column. Rename is nothing but replacement of name and Alias is an alternate name of the table or column.

120) What is a View?

View is a logical table which based on one or more tables or views. The tables upon which the view is based are called Base Tables and it doesn't contain data.

121) What is a cursor variable?

A cursor variable is associated with different statements which can hold different values at run time. A cursor variable is a kind of reference type.

122) What are cursor attributes?

Each cursor in Oracle has set of attributes which enables an application program to test the state of the cursor. The attributes can be used to check whether cursor is opened or closed, found or not found and also find row count.

123) What are SET operators?

SET operators are used with two or more queries and those operators are Union, Union All, Intersect and Minus.

124) How can we delete duplicate rows in a table?

Duplicate rows in the table can be deleted by using ROWID.

125) What are the attributes of Cursor?

Attributes of Cursor are

1. %FOUND

Returns NULL if cursor is open and fetch has not been executed

Returns TRUE if the fetch of cursor is executed successfully.

Returns False if no rows are returned.

- %NOT FOUND

Returns NULL if cursor is open and fetch has not been executed

Returns False if fetch has been executed

Returns True if no row was returned

- %ISOPEN

Returns true if the cursor is open

Returns false if the cursor is closed

- %ROWCOUNT

Returns the number of rows fetched. It has to be iterated through entire cursor to give exact real count.

126) Can we store pictures in the database and if so, how it can be done?

Yes, we can store pictures in the database by Long Raw Data type. This datatype is used to store binary data for 2 gigabytes of length. But the table can have only on Long Raw data type.

127) What is an integrity constraint?

An integrity constraint is a declaration defined a business rule for a table column. Integrity constraints are used to ensure accuracy and consistency of data in a database. There are types – Domain Integrity, Referential Integrity and Domain Integrity.

128) What is an ALERT?

An alert is a window which appears in the center of the screen overlaying a portion of the current display.

129) What is hash cluster?

Hash Cluster is a technique used to store the table for faster retrieval. Apply hash value on the table to retrieve the rows from the table.

130) What are the various constraints used in Oracle?

Following are constraints used:

- NULL – It is to indicate that particular column can contain NULL values
- NOT NULL – It is to indicate that particular column cannot contain NULL values
- CHECK – Validate that values in the given column to meet the specific criteria
- DEFAULT – It is to indicate the value is assigned to default value

131) What is difference between SUBSTR and INSTR?

SUBSTR returns specific portion of a string and INSTR provides character position in which a pattern is found in a string.

SUBSTR returns string whereas INSTR returns numeric.

132) What is the parameter mode that can be passed to a procedure?

IN, OUT and INOUT are the modes of parameters that can be passed to a procedure.

133) What are the different Oracle Database objects?

There are different data objects in Oracle –

- Tables – set of elements organized in vertical and horizontal
- Views – Virtual table derived from one or more tables
- Indexes – Performance tuning method for processing the records
- Synonyms – Alias name for tables
- Sequences – Multiple users generate unique numbers
- Tablespaces – Logical storage unit in Oracle

134) What is the fastest query method to fetch data from the table?

Row can be fetched from table by using ROWID. Using ROW ID is the fastest query method to fetch data from the table.

135) What is the maximum number of triggers that can be applied to a single table?

12 is the maximum number of triggers that can be applied to a single table.

136) How to display row numbers with the records?

Display row numbers with the records numbers – Select rownum, <fieldnames> from table;

This query will display row numbers and the field values from the given table

137) How can we view last record added to a table?

Last record can be added to a table and this can be done by –

Select * from (select * from employees order by rownum desc) where rownum<2;

138) What is the data type of DUAL table?

The **DUAL** table is a one-column table present in oracle database. The table has a single VARCHAR2(1) column called DUMMY which has a value of 'X'.

139) What is difference between Cartesian Join and Cross Join?

There are no differences between the join. Cartesian and Cross joins are same. Cross join gives cartesian product of two tables – Rows from first table is multiplied with another table which is called cartesian product.

Cross join without where clause gives Cartesian product.

140) How to display employee records who gets more salary than the average salary in the department?

This can be done by this query –

Select * from employee where salary>(select avg(salary) from dept, employee where dept.deptno = employee.deptno;

Summary Feature Comparison

The following table includes information about the Oracle, MySQL, and SQL Server databases and how they compare.

Feature	Oracle	MySQL	SQL Server
Interface	GUI, SQL	SQL	GUI, SQL, Various
Language support	Many, including C, C#, C++, Java, Ruby, and Objective C	Many, including C, C#, C++, D, Java, Ruby, and Objective C	Java, Ruby, Python, VB, .Net, and PHP
Operating System	Windows, Linux, Solaris, HP-UX, OS X, z/OS, AIX	Windows, Linux, OS X, FreeBSD, Solaris	Windows
Licensing	Proprietary	Open source	Proprietary
